

CSE322 Computer Network

NS2 Term Project

# Implementation and Evaluation of TCP HyStart in NS2

[Taming the elephants: New TCP slow start - ScienceDirect](#)

Static 802.11

Mobile 802.15.4

HAZ Sameen Shahgir

1805053

Group A2

## Motivation:

When operating under large bandwidth-delay product (BDP) networks, the conventional slow start technique proves to be ineffective. Standard TCP doubles the congestion window (cwnd) for every round-trip time (RTT) during slow start. However, the exponential growth of cwnd results in burst packet losses. Since the cwnd overshoots the path capacity an amount as large as the entire BDP, in large BDP networks, this overshoot causes strong disruption in the networks.

The selective acknowledgement (SACK) option relieves this problem to some extent. As SACK informs the sender the blocks of packets successfully received, the sender can be more intelligent about recovering from multiple packet losses. However, for a large BDP network where a large number of packets are in flight, the processing overhead of SACK information at the end points can be overwhelming. This is because each SACK block invokes a search into the large packet buffers of the sender for the acknowledged packets in the block, and every recovery of a lost packet causes the same search at the receiver end. During fast recovery, every packet reception generates a new SACK packet. Given that the size of the cwnd can be quite large (sometimes, greater than 100,000), the overhead of such a search can be overwhelming.

As such, it is necessary to find a better alternative to the usual technique of waiting for a packet loss to exit slow start. Hystart is an addon to the typical slow start algorithm that exits slow start based on two heuristics, namely, how close the sum of the time intervals between acknowledgements is to the minimum delay time and on whether RTT is abruptly increased or not.

# Description of Hystart Algorithm:

## First Heuristic

### ACK Train Detection

Let the unused available bandwidth of the forward path, the minimum forward path one-way delay and available buffer space of the forward path be  $B$ ,  $D_{min}$  and  $S$  respectively. Then a safe exit point must be less than:

$$C = B \times D_{min} + S$$

$$C < B \times D_{min} = BDP$$

$$\Delta N = \sum_1^N t_k, \text{ where } t = \text{time internal between packets,}$$

$N = \text{Number of packets}$

$$\text{estimated BDP} = b(N) = \frac{(N - 1) \times L}{\Delta N}$$

$$d_{min} = \text{estimated } D_{min}$$

$$\text{estimated } C = C^* = b(N) \times d_{min}$$

$$C^* = \frac{(N - 1) \times L}{\Delta N} \times d_{min}$$

When the upperbound of the network has been reached:  $(N - 1) \times L \rightarrow C^*$   
 $\approx BDP$

then  $\Delta N \rightarrow d_{min}$

## Second Heuristic

### Abrupt Increase in RTT

The first  $n$  packets/their their ACK counterparts are sampled. If  $RTT_k = RTT_{k-1} + \eta$  (where  $\eta$  is a constant), then another safe exit point has been found.

## Code Modification:

# tcp\_cubic.c

ns2-2.35/tcp/linux.src

## My Modifications over TCP Hystart

1. `hystart_ack_delta_us` was doubled (to 4000us) in order to trigger Hystart more often. This parameter controls the upper bound on the time difference between two consecutive ACK such that they are considered a part of an ACK train.
2. Current RTT is not bound by smoothed RTT (constant 100ms) since the 2<sup>nd</sup> heuristic checks whether the Current RTT is greater than Last RTT + constant

```

1 static void hystart(struct sock *sk)
2 {
3     struct tcp_sock *tp = tcp_sk(sk);
4     struct bictcp *ca = inet_csk_ca(sk);
5     u32 threshold;
6     u32 now = bictcp_clock_us(sk);
7
8     if (after(tp->snd_una, ca->end_seq)) {
9         hystart_reset(sk);
10        return;
11    }
12
13    measure_delay(sk); //recalc ca->delay_min
14
15    if ((s32)(now - ca->last_ack) <= hystart_ack_delta_us) {
16        ca->last_ack = now;
17
18
19        threshold = ca->delay_min >> 1;
20
21
22        if ((s32)(now - ca->round_start) > threshold) {
23            printf("heul\n");
24            ca->found = 1;
25            tp->snd_ssthresh = tp->snd_cwnd;
26            return;
27        }
28    }
29
30    if (ca->sample_cnt < HYSTART_MIN_SAMPLES) {
31        ca->sample_cnt++;
32    }
33    else {
34
35        if (ca->curr_rtt > ca->delay_min +
36            HYSTART_DELAY_THRESH(ca->delay_min >> 3)) {
37            printf("hue2\n");
38            ca->found = 1;
39            tp->snd_ssthresh = tp->snd_cwnd;
40        }
41    }
42 }

```

```

1 static inline u64 bictcp_clock_us(const struct sock *sk)
2 {
3     return ktime_get_real/1000;
4 }
5
6 static inline void hystart_reset(struct sock *sk)
7 {
8     struct tcp_sock *tp = tcp_sk(sk);
9     struct bictcp *ca = inet_csk_ca(sk);
10    ca->last_ack = bictcp_clock_us(sk);
11    ca->round_start = bictcp_clock_us(sk);
12    ca->end_seq = tp->snd_nxt;
13    ca->curr_rtt = ~0U;
14    ca->sample_cnt = 0;
15 }

```

```

1 static void bictcp_acked(struct sock *sk, u32 cnt,
2                          ktime_t last)
3 {
4     ...
5
6     // change 1
7     if (!ca->found &&
8         tp->snd_cwnd <= tp->snd_ssthresh){
9         hystart(sk);
10    }
11
12 }

```

**Topology Setup:**

# Static

## MAC Type: 802.11

DSDV

Single Random Source, Multiple Sink

Area: 500x500 m

Default Number of Nodes: 40

Default Number of Flows: 20

Default TX Multiplier: 1

Default Packets per Second: 100

### **Measurement Units:**

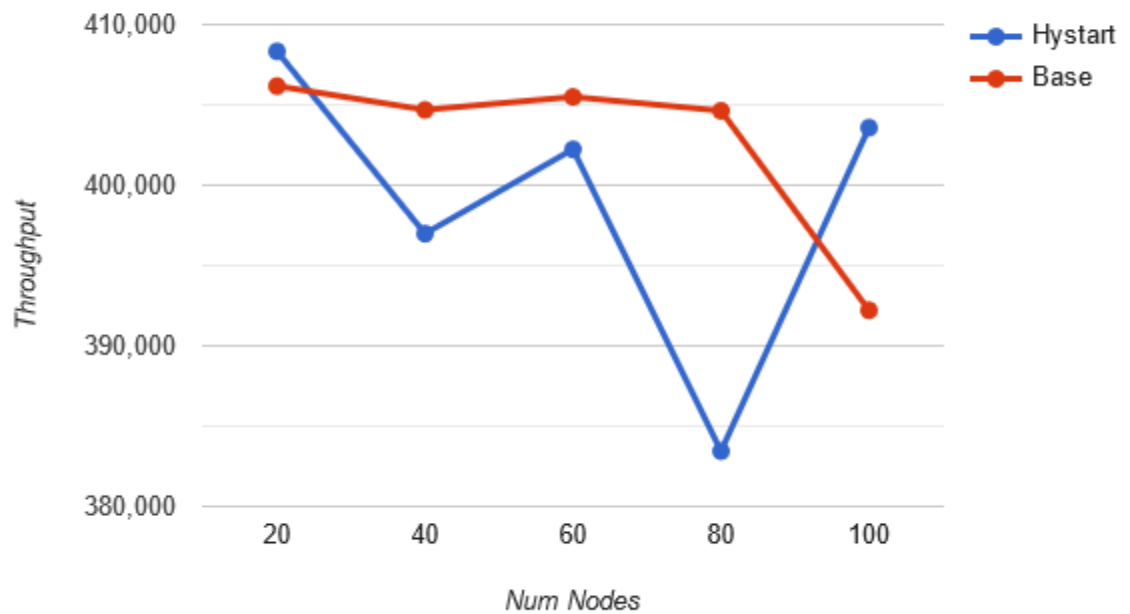
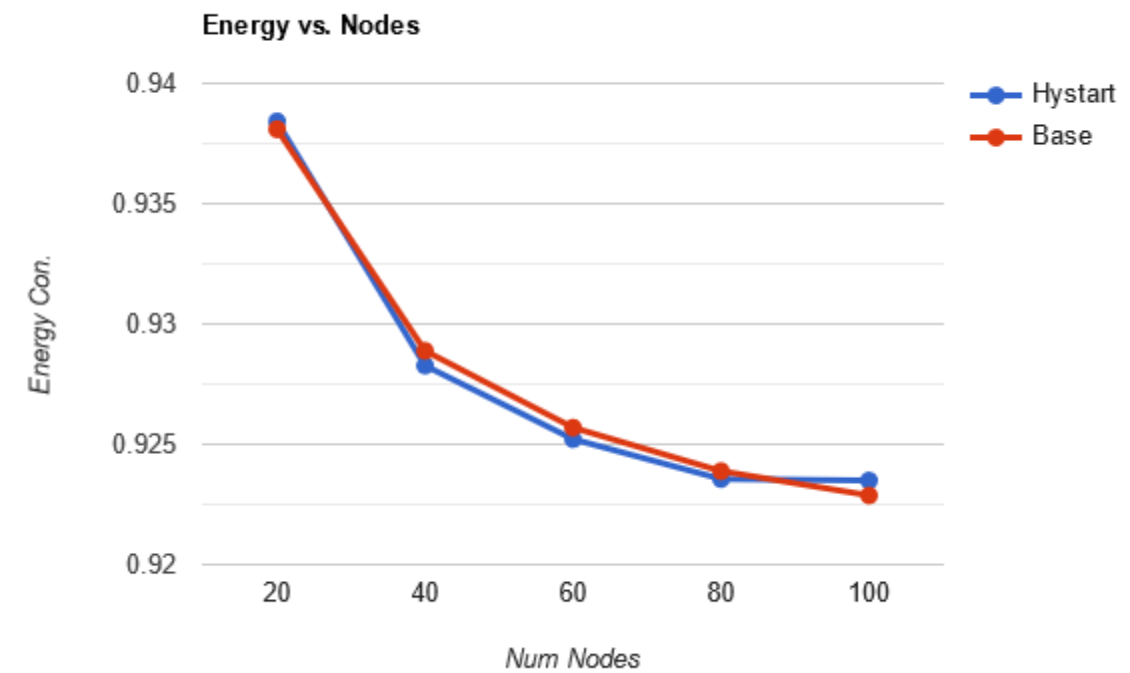
Energy Consumption:  $\text{Js}^{-1}$

Throughput: bps

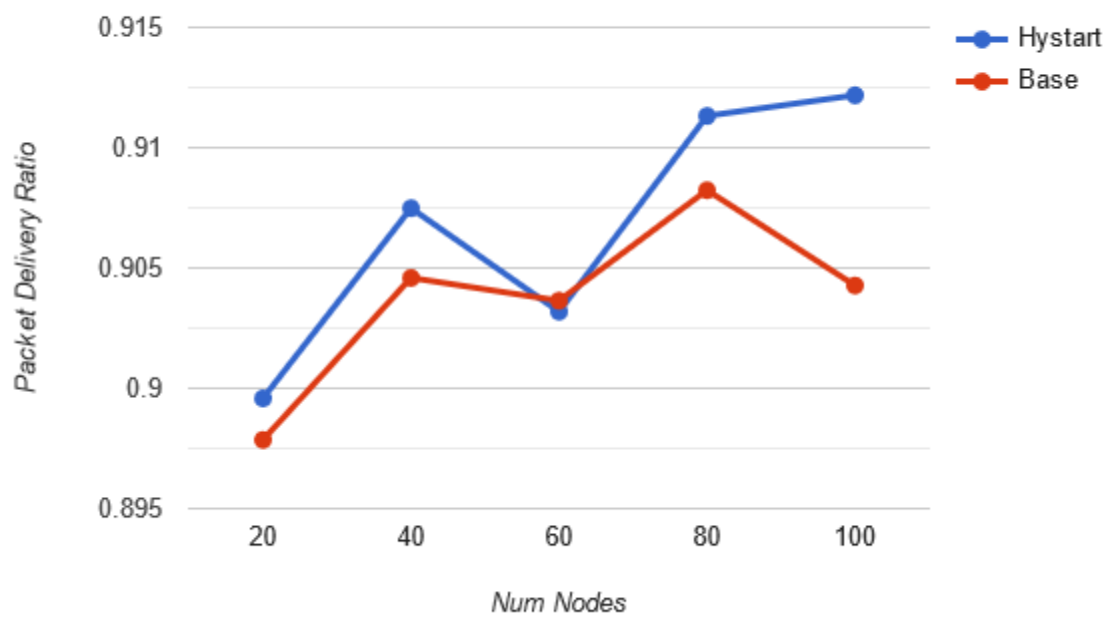
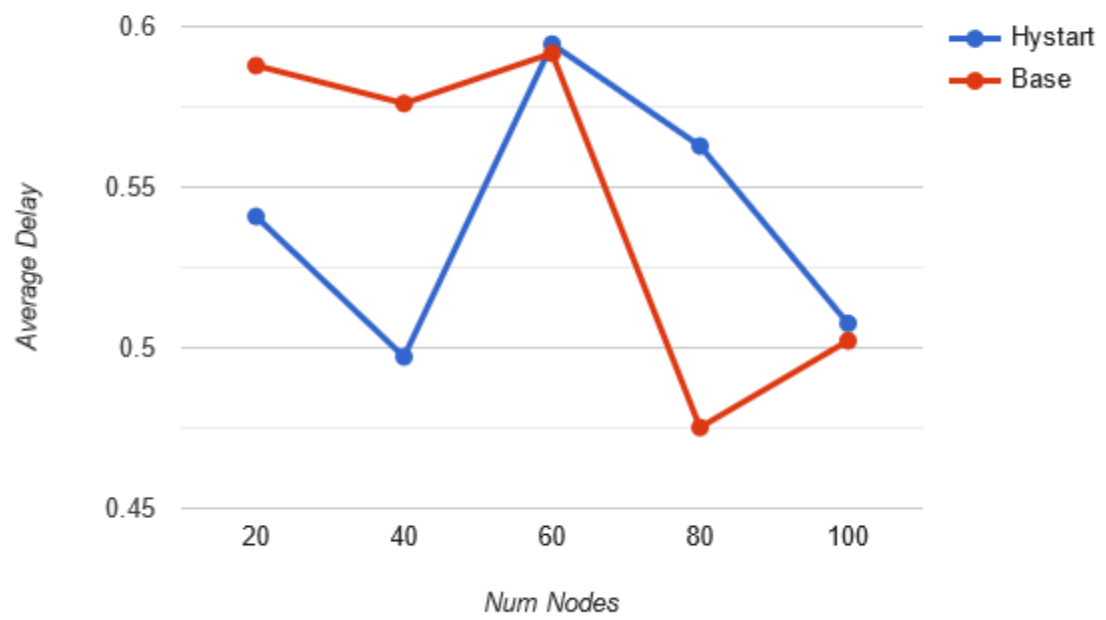
Avg. Delay: s


TX Range: (multiplier)

## 1. Metrics vs. Number of Nodes














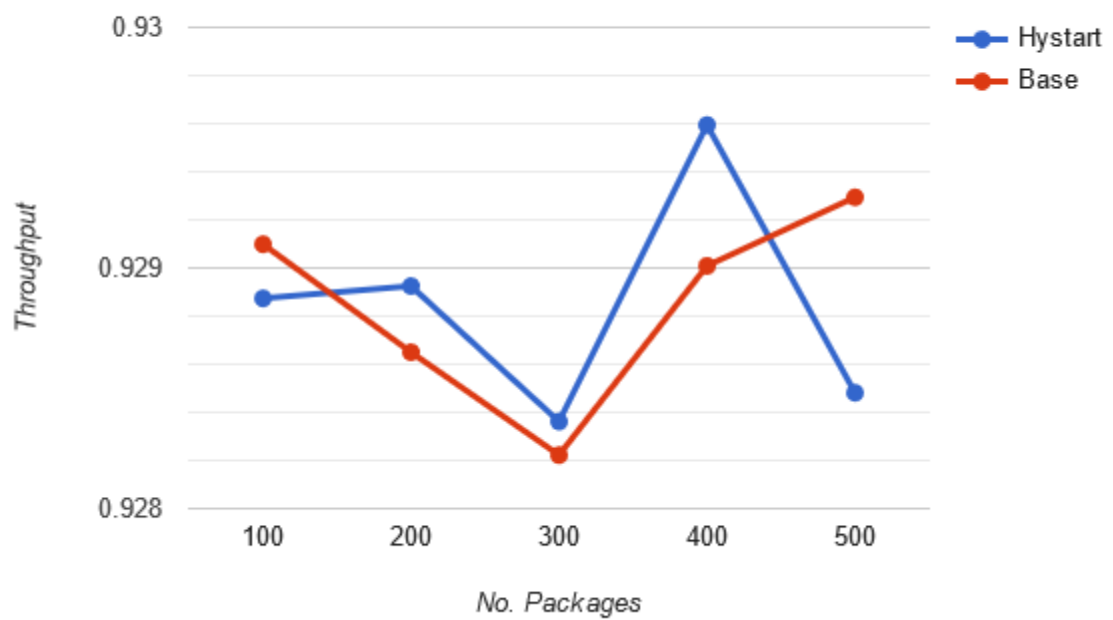
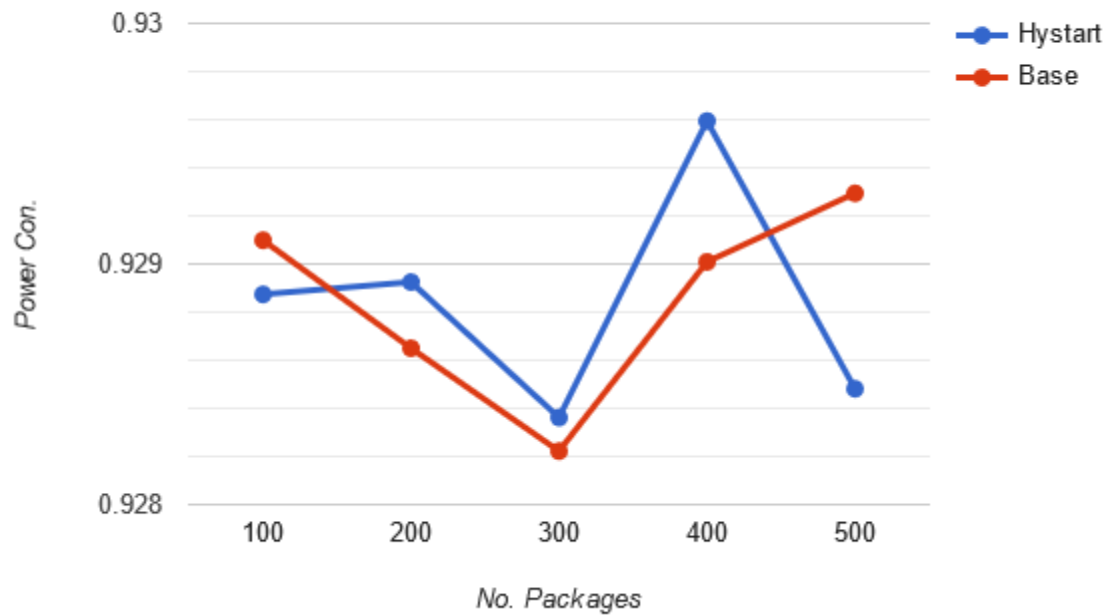
## 2. Metrics vs. Number of Flows

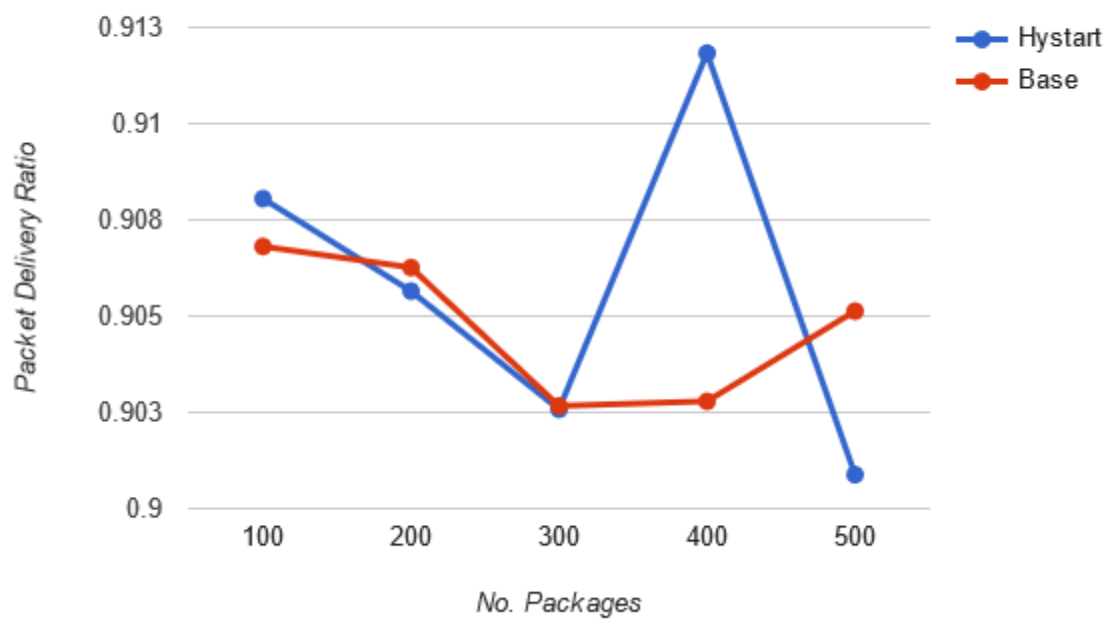
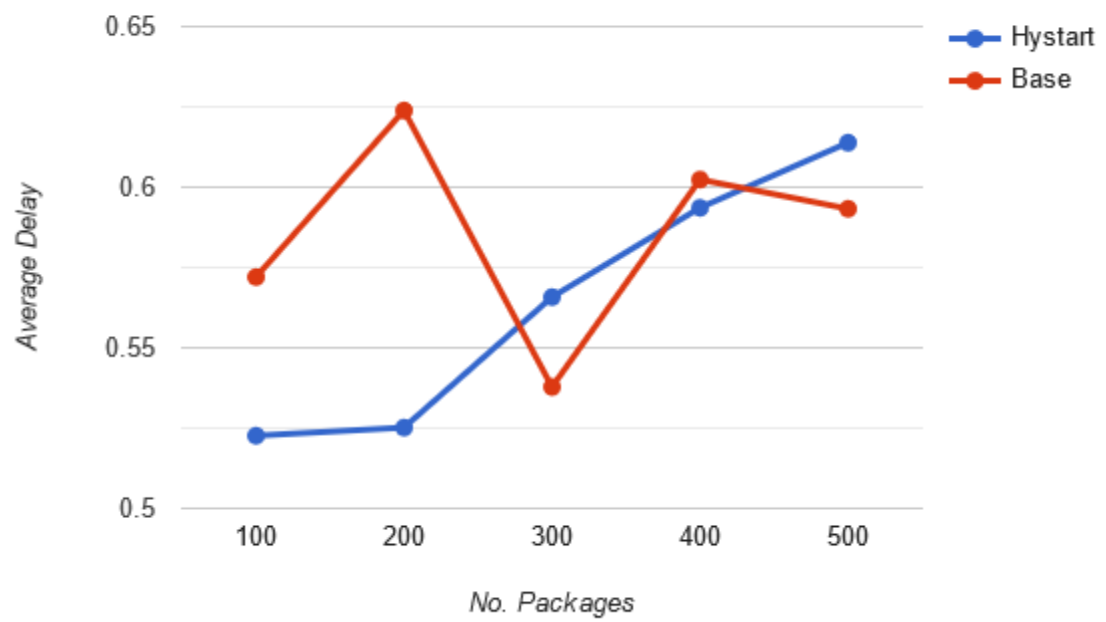


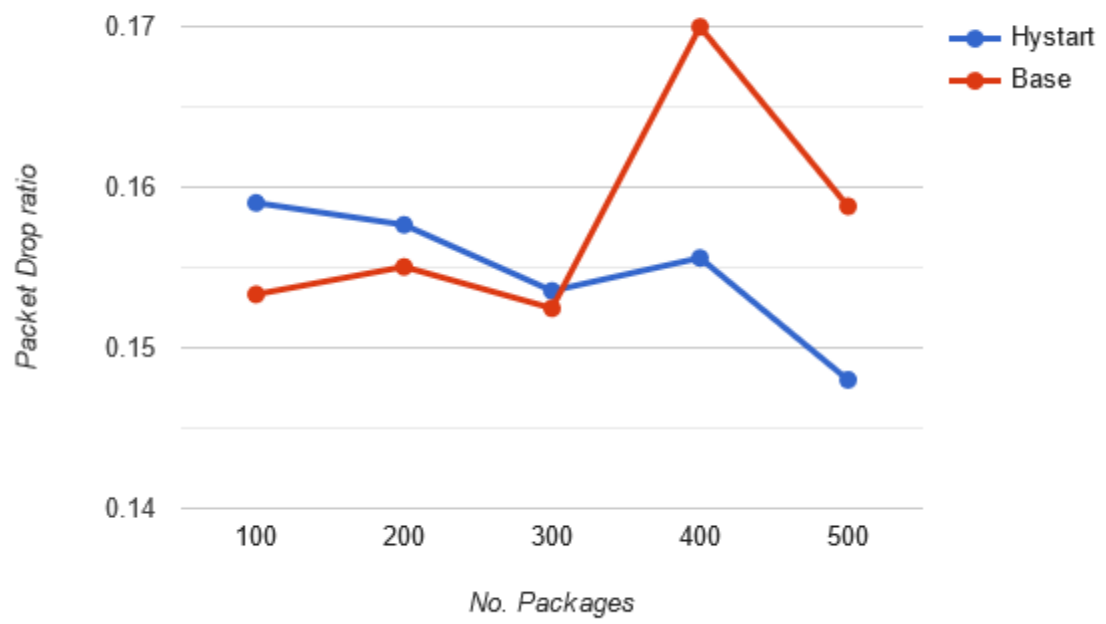




### 3. Metrics vs. Number of Packets per second


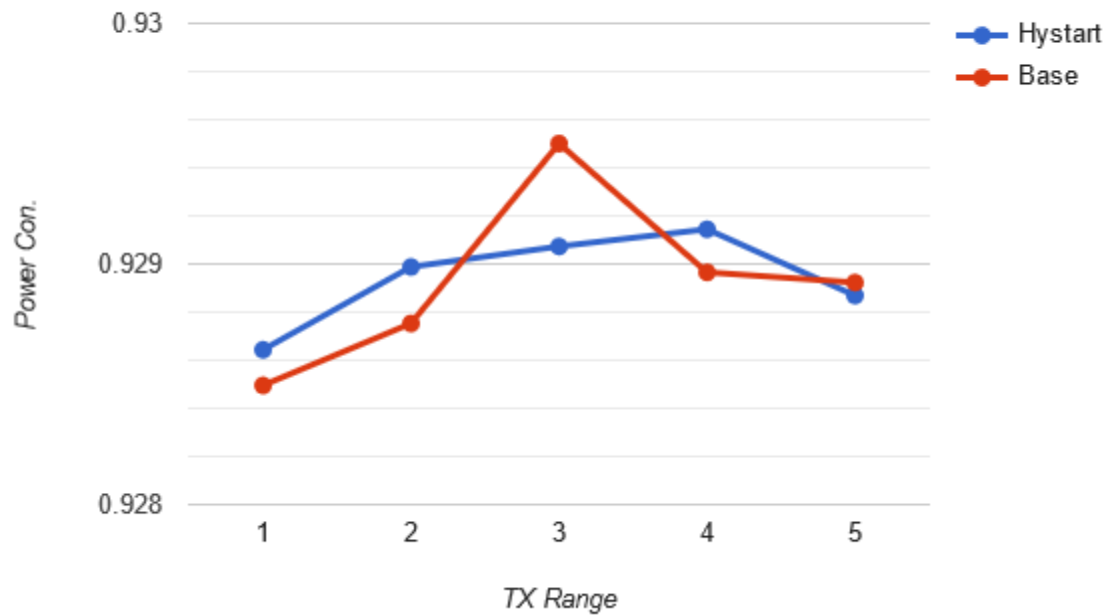


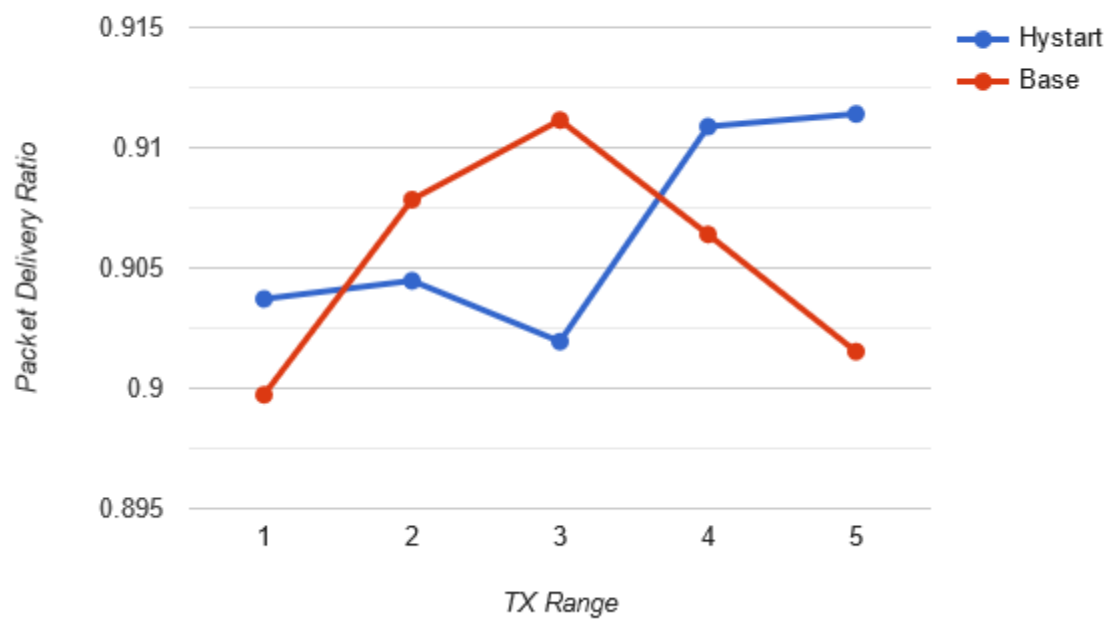
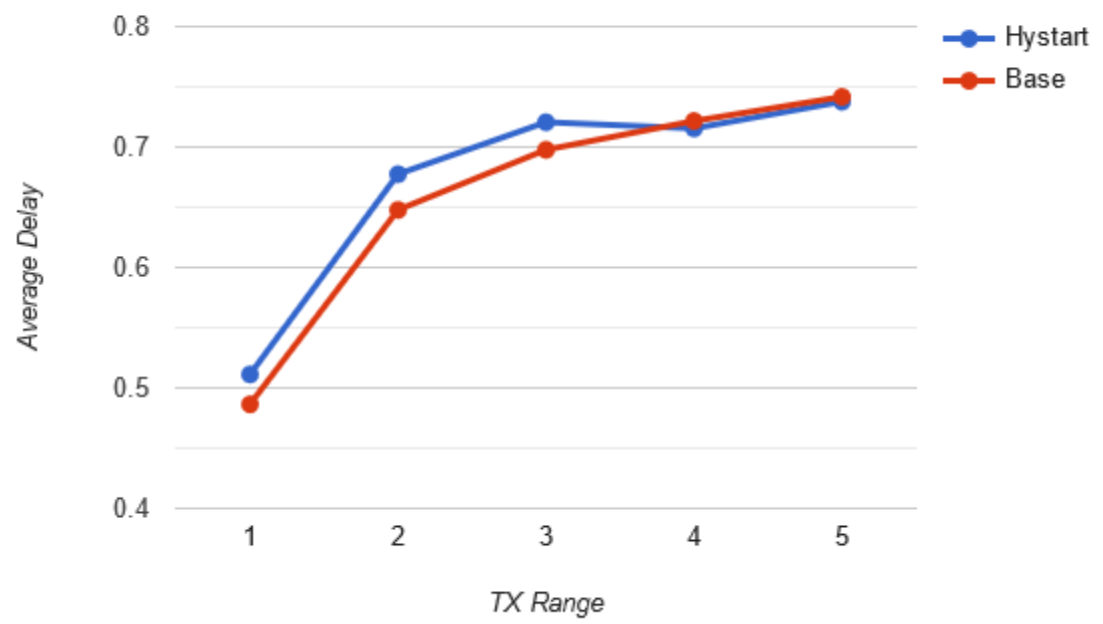







## 4. Metrics vs. TX Range







## **Topology Setup:**

# Mobile

# MAC Type: 802.15.4

DSDV

Single Random Source, Multiple Sink

Area: 500x500 m

Default Number of Nodes: 40

Default Number of Flows: 20

Default Speed: 10

Default Packets per Second: 100

## **Measurement Units:**



Energy Consumption:  $\text{Js}^{-1}$


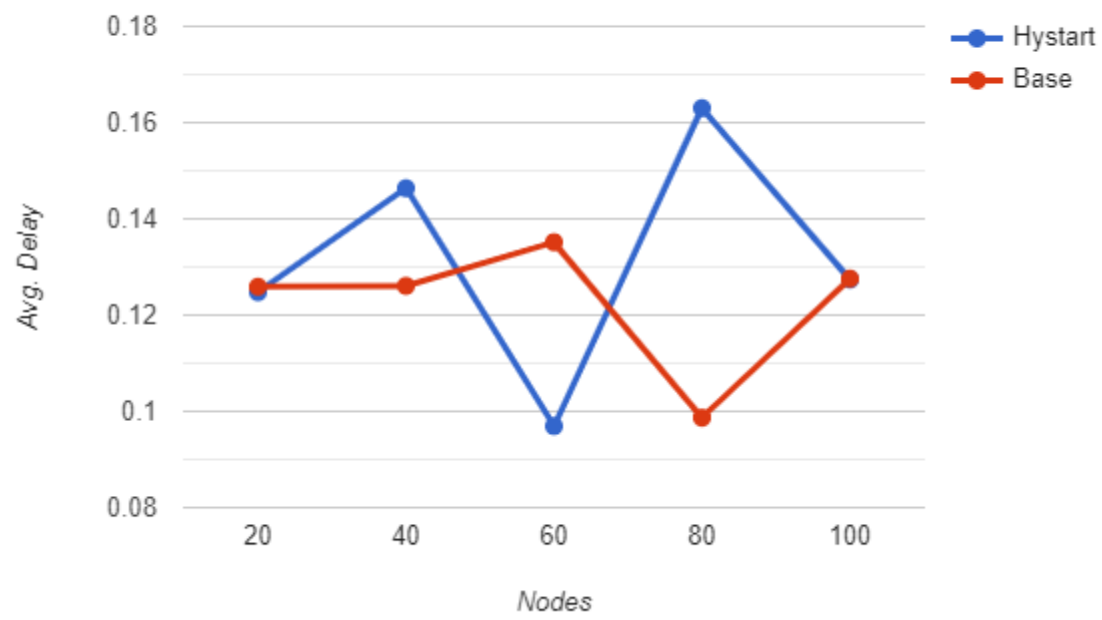
Throughput: bps


Avg. Delay: s

Speed:  $\text{ms}^{-1}$

## 1. Metrics vs. Number of Nodes






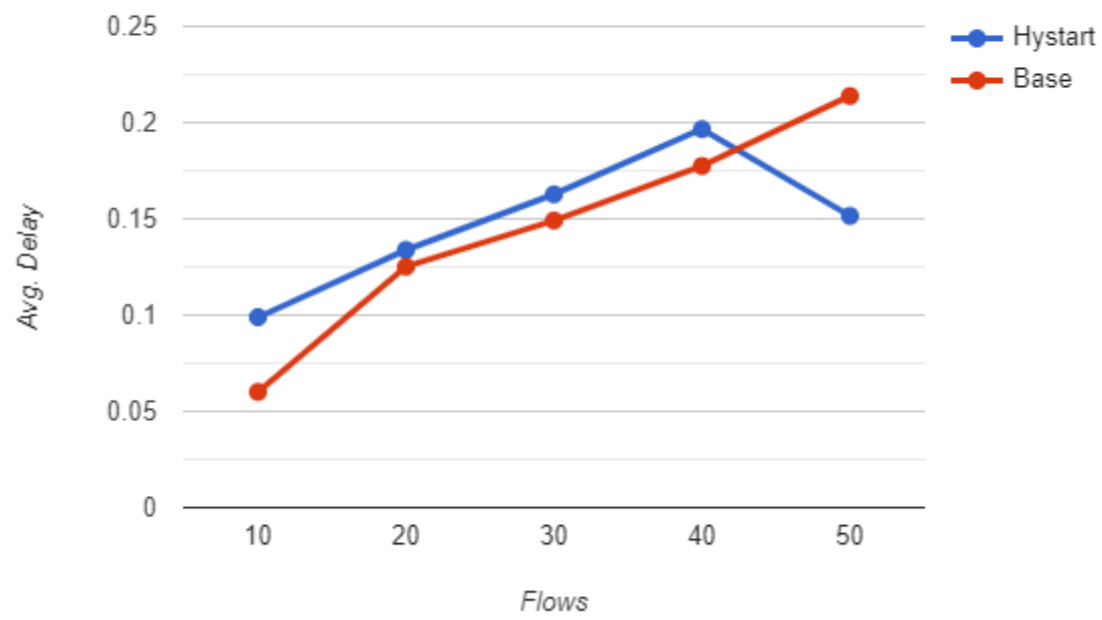


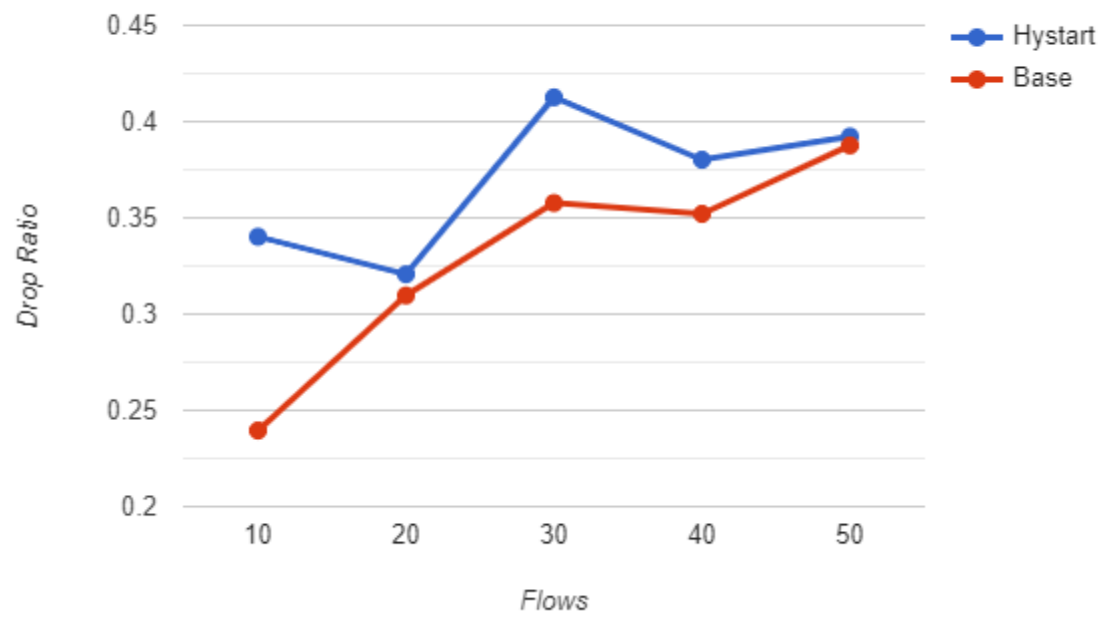





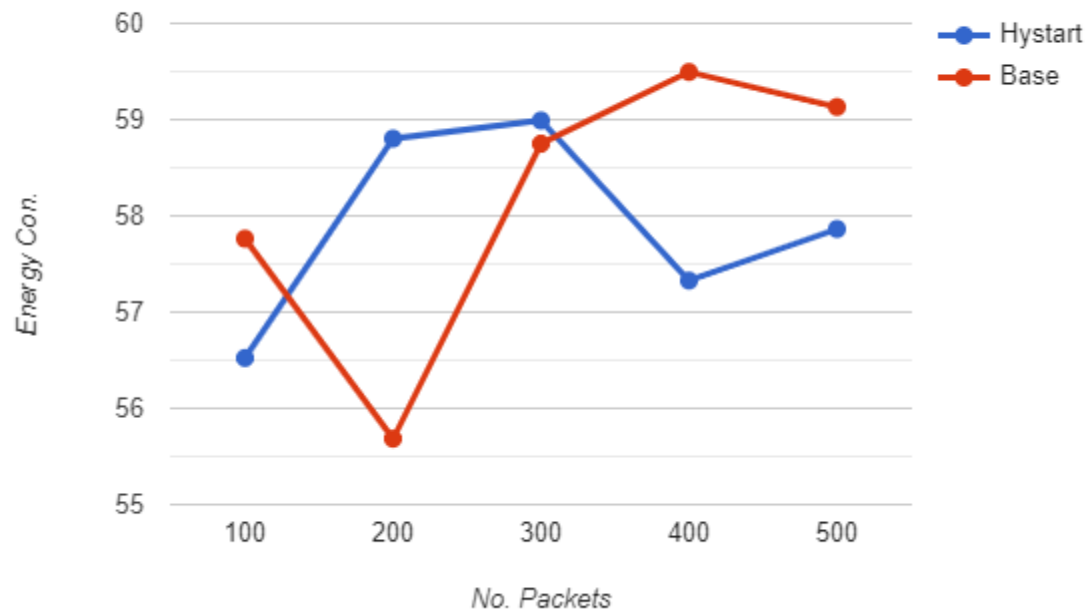
## 2. Metrics vs. Number of Flows












### 3. Metrics vs. Number of Packets Per Second










## 4. Metrics vs. Speed









# Summary Findings:

## In 802.11

1. Energy consumption of TCP Hystart and base TCP Cubic are comparable in all cases.
2. Throughput is slightly but consistently lower in TCP Hystart when measured against Number of Nodes and Number of Flows. In other cases, it's comparable to base TCP Cubic.
3. Average Delay is comparable.
4. Delivery Ratio is improved noticeably when compared.
5. There is a large improvement in Drop Ratio.
6. The comparative results of 50 iterations with the following values:

Iterations = 50  
Area = 500x500 m  
Nodes = 40  
Flow = 20  
MAC = 802.11  
TX Range Multiplier = 3  
Application = FTP  
Packet Size = 1024 bytes  
Packet Interval = 0.005 sec.

| Criteria       | Base       | Hystart    | %Difference    |
|----------------|------------|------------|----------------|
| Throughput     | 418016.244 | 417643.959 | -0.08906       |
| Average Delay  | 0.80988    | 0.798080   | -1.4576        |
| Delivery Ratio | 0.81554    | 0.86891    | <b>6.545</b>   |
| Drop Ratio     | 0.2396     | 0.1906     | <b>-20.434</b> |

## In 802.15.4

Setting MAC type to 802.15.4 in NS2 leads to `packet_length_invalid_error`. To remedy it, the following Energy Mode settings were necessary:

```
set val(initialenergy_15) 300.0
set val(idlepower_15)    40
set val(rxpower_15)      75
set val(txpower_15)      75
set val(sleeppower_15)   40
```

Even still, communication between nodes could not be guaranteed. Due to the random placement of nodes, the nodes were sometimes too far away to communicate. All readings were hence taken with the median of 10 observations, since taking the mean would misrepresent the actual findings.

There seems to be no apparent improvement by TCP Hystart when it comes to 802.15.4. This is likely due to the above-mentioned case where communication via 802.15.4 fails for distant nodes. This also has a knock-on effect on whether Hystart is triggered at all. Since the communication is so sparse, there is seldom an ACK train and hence the Hystart heuristics are rarely triggered. It was observed that a Hystart slow start exit was triggered about 1 in 10 simulations.

# Bonus Work

1. New Measurement:  
**per-node-throughput**
2. Cross-Transmission (Attempted):  
**wired\_wireless.tcl**
3. New Network:  
**Satelite\_iridium.tcl**

# Acknowledgement

Ha, Sangtae, and Injong Rhee. "Taming the elephants: New TCP slow start." *Computer Networks* 55.9 (2011): 2092-2110.